

Week 3: Data Labeling

CS 203: Software Tools and Techniques for AI

Prof. Nipun Batra
IIT Gandhinagar

The Story So Far

Week 1: Collected movie data from OMDb API

Week 2: Validated and cleaned the dataset

Now we have:

- Clean dataset with 50+ movies
- Features: title, year, genre, rating, director, etc.
- Data types validated, missing values handled

This week: Create labeled training data for ML tasks

The reality: Most ML models need labeled data to learn from

What is Data Labeling?

Data labeling: The process of annotating raw data with meaningful labels

Examples:

- Image classification: "This image contains a cat"
- Object detection: Draw boxes around cars in an image
- Sentiment analysis: "This review is positive"
- Named entity recognition: Mark person names in text

Why it matters: Supervised ML models learn from labeled examples

The Netflix Labeling Challenge

New scenario: Netflix wants to build ML models for:

1. **Movie poster classification:** Classify genre from poster images
2. **Review sentiment analysis:** Positive/negative user reviews
3. **Scene detection:** Identify action scenes in trailers
4. **Actor recognition:** Detect which actors appear in scenes

Each task needs labeled training data!

Today's Agenda

1. **Understanding Annotation Tasks** (types and challenges)
2. **Vision Tasks** (classification, detection, segmentation)
3. **Text Tasks** (classification, NER, sentiment)
4. **Label Studio** (annotation tool)
5. **Quality Metrics** (inter-annotator agreement)
6. **Best Practices** (building quality datasets)

Part 1: Understanding Annotation Tasks

What makes a good annotation task?

Types of Machine Learning Tasks

Classification: Assign category labels

- Image: "cat", "dog", "bird"
- Text: "positive", "negative", "neutral"

Detection: Find and locate objects

- Bounding boxes around objects
- Keypoints (facial landmarks)

Segmentation: Pixel-level labels

- Semantic: Label every pixel
- Instance: Separate object instances

Types of ML Tasks (continued)

Sequence Labeling: Tag each element

- Named Entity Recognition (NER)
- Part-of-speech tagging

Structured Prediction: Complex outputs

- Image captioning
- Question answering

Ranking: Order by relevance

- Search results
- Recommendations

Annotation Challenges

1. Subjectivity: Different annotators may disagree

- "Is this movie poster scary or thrilling?"
- Personal interpretation varies

2. Ambiguity: Data can be unclear

- Blurry images
- Ambiguous text

3. Complexity: Some tasks are inherently difficult

- Fine-grained classification (140 dog breeds)
- Multiple objects in one image

Annotation Challenges (continued)

4. Cost: Labeling is expensive

- Time-consuming manual work
- Expert knowledge sometimes required

5. Scale: ML needs large datasets

- ImageNet: 1.2 million labeled images
- Small datasets lead to overfitting

6. Quality: Errors propagate to the model

- Noisy labels hurt performance
- Need quality control measures

Part 2: Vision Annotation Tasks

Labeling images and videos

Image Classification

Task: Assign one label to an entire image

Example - Movie Poster Genre:

Image: [poster of Inception]

Label: Sci-Fi

Use cases:

- Content moderation
- Medical diagnosis
- Product categorization

Annotation: Simple - just pick a category

Multi-Label Classification

Task: Assign multiple labels to one image

Example - Movie Poster Genres:

Image: [poster of Inception]
Labels: ["Sci-Fi", "Thriller", "Action"]

Difference from single-label:

- One image can have multiple categories
- More flexible but more complex

Object Detection

Task: Find and locate objects with bounding boxes

Example - Actor Detection:

Image: [movie scene]

Boxes:

- [x:100, y:50, w:80, h:120] "Leonardo DiCaprio"
 - [x:250, y:60, w:75, h:115] "Tom Hardy"

Output: Box coordinates + class label for each object

Use cases: Self-driving cars, face detection, security

Bounding Box Format

Common formats:

1. [x, y, width, height]: Top-left corner + dimensions

```
[100, 50, 80, 120]
```

2. [x1, y1, x2, y2]: Top-left + bottom-right corners

```
[100, 50, 180, 170]
```

3. [x_center, y_center, width, height]: YOLO format

```
[140, 110, 80, 120]
```

Semantic Segmentation

Task: Label every pixel with a class

Example - Scene Segmentation:

```
Pixel [0,0]: sky  
Pixel [0,1]: sky  
Pixel [100,200]: building  
Pixel [150,300]: road
```

Use cases: Medical imaging, autonomous driving, satellite imagery

Annotation: Time-consuming - must trace object boundaries

Instance Segmentation

Task: Separate different instances of same class

Example - Multiple Actors:

Mask 1 (Person): Actor A's pixels
Mask 2 (Person): Actor B's pixels

Difference from semantic:

- Semantic: All actors labeled as "person"
- Instance: Each actor gets separate mask

More precise but more expensive to annotate

Keypoint Detection

Task: Mark specific points of interest

Example - Facial Landmarks:

Points:

- Left eye: (150, 100)
- Right eye: (200, 100)
- Nose: (175, 130)
- Mouth corners: (160, 160), (190, 160)

Use cases: Pose estimation, facial recognition, gesture control

Part 3: Text Annotation Tasks

Labeling text data

Text Classification

Task: Assign category to entire text

Example - Review Sentiment:

Text: "This movie was absolutely amazing!"

Label: Positive

Use cases:

- Spam detection
- Topic classification
- Sentiment analysis

Annotation: Read text, pick category

Named Entity Recognition (NER)

Task: Identify and classify entities in text

Example - Movie Review:

Text: "Christopher Nolan directed Inception in 2010."

Entities:

- "Christopher Nolan" [PERSON]
 - "Inception" [MOVIE]
 - "2010" [DATE]

Common entity types: PERSON, ORG, LOC, DATE, MONEY

NER Annotation Format

BIO Tagging (Beginning, Inside, Outside):

Christopher	B-PERSON
Nolan	I-PERSON
directed	O
Inception	B-MOVIE
in	O
2010	B-DATE
.	O

B-: Beginning of entity

I-: Inside entity (continuation)

O: Outside any entity

Span-Based NER

Alternative format: Character offsets

Text: "Christopher Nolan directed Inception"

Spans:

- [0, 17]: PERSON
- [27, 36]: MOVIE

Advantages:

- Handles nested entities
- Easier for some tools

Relation Extraction

Task: Identify relationships between entities

Example:

Text: "Christopher Nolan directed Inception."

Relations:

- (Christopher Nolan, directed, Inception)
 - Type: DIRECTOR_OF

Use cases: Knowledge graphs, question answering

Text Sequence Labeling

Task: Label each token

Example - Part-of-Speech Tagging:

The	DET
movie	NOUN
was	VERB
great	ADJ
.	PUNCT

Other examples:

- Chunking (noun phrases, verb phrases)
- Slot filling (book flight to NYC on Monday)

Part 4: Label Studio

Open-source annotation tool

What is Label Studio?

Label Studio: Web-based annotation platform

Features:

- Multi-modal: Images, text, audio, video
- Customizable: Configure for any task
- Collaborative: Multiple annotators
- Export: ML-ready formats

Why use it?:

- Free and open-source
- Easy to set up
- Professional quality

Installing Label Studio

Option 1: pip install (recommended):

```
pip install label-studio  
label-studio start
```

Option 2: Docker:

```
docker run -p 8080:8080 heartexlabs/label-studio
```

Access: Open browser to <http://localhost:8080>

Label Studio Interface

Main components:

1. **Projects:** Collection of annotation tasks
2. **Data:** Import your images/text
3. **Labeling Config:** Define annotation schema
4. **Annotators:** Assign work to people
5. **Export:** Download labeled data

Workflow: Create project → Import data → Configure labels → Annotate → Export

Creating a Project

Step 1: Click "Create Project"

Step 2: Give it a name ("Movie Poster Classification")

Step 3: Import data (upload images or provide URLs)

Step 4: Set up labeling interface

Step 5: Start annotating

Labeling Config: Image Classification

XML-based configuration:

```
<View>
  <Image name="image" value="$image"/>
  <Choices name="genre" toName="image">
    <Choice value="Action"/>
    <Choice value="Comedy"/>
    <Choice value="Drama"/>
    <Choice value="Horror"/>
    <Choice value="Sci-Fi"/>
    <Choice value="Romance"/>
  </Choices>
</View>
```

This creates: Image viewer + genre selector

Labeling Config: Object Detection

For bounding boxes:

```
<View>
  <Image name="image" value="$image"/>
  <RectangleLabels name="label" toName="image">
    <Label value="Actor"/>
    <Label value="Text"/>
    <Label value="Logo"/>
  </RectangleLabels>
</View>
```

This creates: Image viewer + draw rectangle tool

Labeling Config: Text Classification

For sentiment analysis:

```
<View>
  <Text name="text" value="$text"/>
  <Choices name="sentiment" toName="text">
    <Choice value="Positive"/>
    <Choice value="Negative"/>
    <Choice value="Neutral"/>
  </Choices>
</View>
```

This creates: Text display + sentiment selector

Labeling Config: NER

For named entity recognition:

```
<View>
  <Text name="text" value="$text"/>
  <Labels name="label" toName="text">
    <Label value="PERSON"/>
    <Label value="MOVIE"/>
    <Label value="DATE"/>
    <Label value="ORGANIZATION"/>
  </Labels>
</View>
```

This creates: Text viewer + highlight tool for entities

Importing Data

Format 1: JSON:

```
[  
  {  
    "image": "https://example.com/poster1.jpg",  
    "title": "Inception"  
  },  
  {  
    "image": "https://example.com/poster2.jpg",  
    "title": "The Matrix"  
  }  
]
```

Format 2: CSV:

image	title
poster1.jpg	Inception
poster2.jpg	The Matrix

Annotation Workflow

- 1. Review task:** See the image/text
- 2. Apply labels:** Click/select/draw
- 3. Submit:** Save annotation
- 4. Next task:** Move to next item

Keyboard shortcuts speed up annotation:

- Number keys: Select label
- Enter: Submit
- Skip: Skip difficult examples

Managing Annotators

Roles:

- **Annotator:** Labels data
- **Reviewer:** Checks quality
- **Manager:** Assigns tasks

Assignment strategies:

- Round-robin: Equal distribution
- Overlap: Multiple people label same item
- Expert: Hard cases go to best annotators

Exporting Labeled Data

Format options:

1. JSON: Full detail

```
{  
  "task": 1,  
  "result": [  
    {"value": {"choices": ["Sci-Fi"]},  
     "from_name": "genre"}]  
}
```

2. CSV: Simplified

image	genre
poster1.jpg	Sci-Fi
poster2.jpg	Action

Export Formats (continued)

3. COCO (for object detection):

```
{  
  "images": [...],  
  "annotations": [{  
    "image_id": 1,  
    "category_id": 1,  
    "bbox": [100, 50, 80, 120]  
  }],  
  "categories": [...]  
}
```

4. YOLO (for detection):

```
0 0.5 0.5 0.2 0.3  
1 0.7 0.6 0.15 0.25
```

Part 5: Quality Metrics

Measuring annotation quality

Why Measure Agreement?

Problem: Different annotators may disagree

Example:

Movie review: "The movie was okay."

Annotator A: Positive
Annotator B: Neutral
Annotator C: Negative

Need to measure: How much do annotators agree?

Goal: High agreement = clear task + good annotators

Simple Agreement

Percent agreement: How often annotators agree

Formula:

$$\text{Agreement} = (\# \text{ agreements}) / (\# \text{ total items})$$

Example:

100 items, both annotators agree on 85

$$\text{Agreement} = 85 / 100 = 0.85 \text{ (85\%)}$$

Problem: Doesn't account for chance agreement

Cohen's Kappa

Better metric: Accounts for random chance

Formula:

$$\kappa = (p_o - p_e) / (1 - p_e)$$

p_o = observed agreement

p_e = expected agreement by chance

Range: -1 to 1

- $\kappa = 1$: Perfect agreement
- $\kappa = 0$: No agreement beyond chance
- $\kappa < 0$: Less than chance (bad!)

Cohen's Kappa Example

Data: 100 movie posters, 2 annotators

Results:

	Ann B: Action	Ann B: Drama
Ann A: Action	30	10
Ann A: Drama	5	55

Observed agreement: $(30 + 55) / 100 = 0.85$

Expected by chance:

- $P(\text{both say Action}) = 0.40 \times 0.35 = 0.14$
- $P(\text{both say Drama}) = 0.60 \times 0.65 = 0.39$
- $p_e = 0.14 + 0.39 = 0.53$

$$\text{Kappa: } (0.85 - 0.53) / (1 - 0.53) = 0.68$$

Interpreting Kappa

Guidelines (Landis & Koch):

- < 0 : Poor
- $0.00 - 0.20$: Slight
- $0.21 - 0.40$: Fair
- $0.41 - 0.60$: Moderate
- $0.61 - 0.80$: Substantial
- $0.81 - 1.00$: Almost perfect

Our example: $\kappa = 0.68$ (Substantial agreement)

Calculating Kappa in Python

```
from sklearn.metrics import cohen_kappa_score

        # Annotations from two annotators
annotator_a = ['Action', 'Drama', 'Action', 'Drama', ...]
annotator_b = ['Action', 'Drama', 'Drama', 'Drama', ...]

kappa = cohen_kappa_score(annotator_a, annotator_b)
print(f"Cohen's Kappa: {kappa:.3f}")
```

Output: Cohen's Kappa: 0.680

Fleiss' Kappa

For 3+ annotators: Extension of Cohen's Kappa

Use case: Multiple annotators label same data

Formula: More complex, measures agreement across all annotators

Python:

```
from statsmodels.stats.inter_rater import fleiss_kappa

    # Matrix: rows = items, cols = categories
    # Values = count of annotators who chose category
    data = [
        [2, 1, 0],  # Item 1: 2 said Action, 1 said Drama
        [0, 3, 0],  # Item 2: 3 said Drama
        ...
    ]
    kappa = fleiss_kappa(data)
```

Confusion Matrix

Visualize disagreements:

		Annotator B		
		Action	Drama	Comedy
Annotator A:				
Action		30	8	2
Drama		5	45	10
Comedy		2	7	21

Diagonal: Agreements

Off-diagonal: Disagreements

Insights: Where do annotators disagree most?

Improving Agreement

Strategies:

1. Better guidelines:

- Define categories clearly
- Provide examples
- Handle edge cases

2. Training:

- Practice sessions
- Discuss disagreements
- Calibrate understanding

3. Task simplification:

Improving Agreement (continued)

4. Quality control:

- Gold standard examples
- Regular agreement checks
- Feedback to annotators

5. Redundancy:

- Multiple annotations per item
- Majority vote or expert resolution
- Identify difficult examples

Part 6: Best Practices

Building high-quality labeled datasets

Annotation Guidelines

Essential components:

1. **Task description:** What are we labeling?
2. **Label definitions:** Clear definition of each category
3. **Examples:** Show good and bad annotations
4. **Edge cases:** How to handle ambiguous cases
5. **Process:** Step-by-step instructions

Document everything!

Example Guidelines: Sentiment

Task: Label movie reviews as Positive, Negative, or Neutral

Definitions:

- **Positive:** Overall favorable opinion
- **Negative:** Overall unfavorable opinion
- **Neutral:** Mixed feelings or factual statement

Examples:

- "Amazing movie!" → Positive
- "Terrible waste of time" → Negative
- "The movie was 2 hours long" → Neutral

Edge cases:

Pilot Annotation

Before full-scale labeling:

1. **Small pilot (10-20 examples)**
2. **Multiple annotators label same data**
3. **Measure agreement**
4. **Identify issues:**
 - Unclear guidelines
 - Ambiguous categories
 - Missing edge cases
5. **Revise and repeat**

Don't skip this step!

Quality Control Strategies

1. Gold standard:

- Expert-labeled examples
- Test annotators periodically
- Track performance

2. Redundancy:

- 2-3 annotations per item
- Resolve disagreements
- Filter low-quality work

3. Regular reviews:

- Spot-check annotations

Handling Disagreements

When annotators disagree:

Option 1: Majority vote

Item 1: A=Positive, B=Positive, C=Negative
Label: Positive (2/3)

Option 2: Expert adjudication

- Expert reviews disagreements
- Makes final decision

Option 3: Discard

- If no clear answer, skip item
- Keep only high-confidence labels

Dataset Split Strategy

Standard splits:

Training (70-80%): Model learns from this

Validation (10-15%): Tune hyperparameters

Test (10-15%): Final evaluation

Important: Keep annotators consistent within splits!

Bad: Different annotators for train vs test

Good: Same quality standards throughout

Annotation Cost Estimation

Factors:

- Task complexity (simple classification vs detailed segmentation)
- Annotator expertise (crowd vs expert)
- Quality requirements (single vs multiple annotations)
- Dataset size (hundreds vs millions)

Example estimates:

- Image classification: \$0.01 - \$0.10 per image
- Bounding boxes: \$0.10 - \$1.00 per image
- Segmentation: \$1.00 - \$10.00 per image
- NER: \$0.05 - \$0.50 per sentence

Active Learning

Reduce annotation cost:

Strategy:

1. Label small initial dataset
2. Train initial model
3. **Model selects** most useful examples to label next
4. Label those examples
5. Retrain and repeat

Key idea: Label the data that helps the model most

Savings: Can reduce labeling by 50-90%

Common Pitfalls

- 1. Vague guidelines:** Leads to disagreement
- 2. Too many classes:** Hard to distinguish
- 3. Imbalanced data:** Too few examples of some classes
- 4. Batch effects:** Annotators change behavior over time
- 5. No validation:** Don't know if labels are good
- 6. Ignoring context:** Labels without understanding domain

Solution: Plan carefully, test early, iterate

Summary: Labeling Workflow

Complete process:

1. **Define task:** What are we predicting?
2. **Create guidelines:** Clear definitions + examples
3. **Pilot test:** Small batch, measure agreement
4. **Revise:** Fix issues identified
5. **Scale up:** Full annotation with quality control
6. **Validate:** Check agreement throughout
7. **Export:** ML-ready format
8. **Split:** Train/val/test sets

Result: High-quality labeled dataset for ML

Next Steps

Homework:

- Set up Label Studio
- Create annotation project for your Netflix data
- Annotate sample of movie posters or reviews
- Measure inter-annotator agreement

Next topics:

- Training models on labeled data
- Evaluation metrics
- Active learning strategies

Key Takeaways

1. **Labeling is critical:** Quality labels = quality models
2. **Many task types:** Classification, detection, NER, etc.
3. **Tools help:** Label Studio makes annotation easier
4. **Measure quality:** Use Kappa to check agreement
5. **Guidelines matter:** Clear instructions improve consistency
6. **Iterate:** Pilot → revise → scale up

Resources

Tools:

- Label Studio: <https://labelstud.io/>
- CVAT (computer vision): <https://cvat.org/>
- Prodigy (text): <https://prodi.gy/>

Metrics:

- scikit-learn metrics: https://scikit-learn.org/stable/modules/model_evaluation.html
- Agreement calculators

Reading:

- Best practices for data annotation
- Active learning tutorials

Questions?

Next class: Lab session - hands-on annotation with Label Studio