

# **Week 10 Lab: Building ML APIs with FastAPI**

**CS 203: Software Tools and Techniques for AI**

Duration: 3 hours

# Lab Objectives

- 1. FastAPI Basics:** Build your first Hello World API.
- 2. Data Validation:** Use Pydantic to secure inputs.
- 3. Model Serving:** Wrap a scikit-learn model in a REST API.
- 4. Testing:** Verify your API works correctly.

## Prerequisites:

- `pip install fastapi[standard] scikit-learn joblib pytest httpx`

# Exercise 1: Hello FastAPI (30 min)

**Goal:** Get a server running and explore Swagger UI.

1. Create `main.py`.
2. Define a root endpoint `GET /` returning `{"message": "ML API is running"}`.
3. Define an endpoint `GET /add/{a}/{b}` that returns the sum of two integers.
4. Run with `fastapi dev main.py`.
5. Open `http://127.0.0.1:8000/docs`.
  - Try the "Execute" button for both endpoints.

## Exercise 2: Pydantic Validation (45 min)

**Goal:** Define a robust schema for an ML model input.

Imagine a "House Price Prediction" model.

Inputs: `square_feet` (int), `bedrooms` (int), `zipcode` (str).

1. Define a `HouseFeatures` Pydantic model.

- `square_feet` : Must be  $> 0$ .
- `bedrooms` : Must be  $\geq 0$ .

2. Create a `POST /predict_price` endpoint.

- Accept `HouseFeatures` as JSON body.
- Return a dummy prediction (e.g., `square_feet * 100` ).

**Test:** Send invalid data (negative size) via Swagger UI and observe the 422 Error.

# Exercise 3: Serving a Real Model (60 min)

**Goal:** The core MLOps task.

1. **Train:** Create `train_model.py`.

- Train a `RandomForestClassifier` on the **Iris dataset**.
- Save it using `joblib.dump(model, "iris_model.pkl")`.
- Run this script once.

2. **Serve:** Update `main.py`.

- Load `iris_model.pkl` on startup (global variable).
- Create `POST /predict_species`.
- Input: `sepal_length`, `sepal_width`, etc. (Use Pydantic!).
- Output: The predicted class name (Setosa/Versicolor/Virginica).

# Exercise 4: Testing (30 min)

**Goal:** Automate API verification.

1. Create `test_api.py`.
2. Use `fastapi.testclient.TestClient`.
3. Write a test `test_predict_species()`:
  - Send a sample JSON payload.
  - Assert status code is 200.
  - Assert "class" is in the response body.
4. Run `pytest`.

# Submission

## Deliverables:

1. `main.py` : Your FastAPI application.
2. `train_model.py` : The training script.
3. `test_api.py` : Your tests.
4. `requirements.txt` : Dependencies.

## Challenge (Bonus):

Add a `/health` endpoint that checks if the model is loaded correctly and returns 200 OK or 500 Error.