

# **Model Deployment**

**CS 203: Software Tools and Techniques for AI**

Prof. Nipun Batra, IIT Gandhinagar

# From Training to Production

## The Journey:

1. Train model on laptop/server
2. Evaluate performance
3. Package model
4. Deploy to production
5. Serve predictions
6. Monitor performance
7. Update and retrain

## Key Concept: The "Wall of Confusion"

Developers write code, Ops deploy it. In MLOps, Data Scientists train models, but Engineers deploy them.

# Deployment Options Overview

Strategy	Latency	Throughput	Use Case
REST API	Low (ms)	High	Real-time (Chatbots, Recommendations)
Batch	High (hrs)	Very High	Nightly Reports, Churn Prediction
Edge	Very Low	Low	IoT, Privacy-sensitive (FacID)
Streaming	Low (ms)	High	Fraud Detection, Sensor Data

```
graph TD A[Client Request] --> B{Type?}; B -- Online --> C[REST API]; B -- Offline --> D[Batch Job]; B -- Event --> E[Stream Processor];
```

# Model Quantization: Theory

**Why?** Models are huge.

- ResNet-50: ~98MB (Float32)
- Quantized: ~25MB (Int8) -> **4x smaller, 2-4x faster**

**How it works:** Map continuous float values to discrete integers.

$$Q(x) = \text{round}(x/S + Z)$$

- $x$ : Input float
- $S$ : Scale factor
- $Z$ : Zero point
- $Q(x)$ : Output integer

**Trade-off:** Precision vs. Size.

Usually accuracy drop is < 1% for standard models

# Quantization Visualized

Float32 has a huge dynamic range. Int8 has only 256 values [-128, 127].

```
graph LR A[Float32 Range] -- Mapping --> B[Int8 Range]; A --> |-10.5| B1[-128]; A --> |0.0| B2[0]; A --> |10.5| B3[127]; style A fill:#e1f5fe style B fill:#fff9c4
```

**PyTorch Code (Post-Training Static Quantization):**

```
import torch

# 1. Define model
model = MyModel()
model.eval()

# 2. Fuse layers (Conv+BN+ReLU) for speed
model.fuse_model()

# 3. Prepare config (x86 or ARM)
model.qconfig = torch.quantization.get_default_qconfig('fbgemm')
torch.quantization.prepare(model, inplace=True)
```

# ONNX: The Universal Bridge

**Problem:** PyTorch models don't run in TensorFlow. Deployment hardware (NVIDIA, Intel) needs specific optimization.

**Solution:** ONNX (Open Neural Network Exchange).

- A common graph representation.
- Write in *any* framework -> Export to ONNX -> Run on *any* hardware.

```
graph LR A[PyTorch] --> D[ONNX Graph]; B[TensorFlow] --> D; C[Scikit-Learn] --> D; D --> E[ONNX Runtime]; E --> F[CPU]; E --> G[NVIDIA GPU]; E --> H[Mobile (ARM)];
```

# Serving Architecture: Containerization

## Why Docker?

"It works on my machine" is not a deployment strategy.

Containers package code + dependencies + OS libraries.

## Comparing Virtual Machines vs Containers:

### Virtual Machine

- Heavy (GBs)
- Guest OS per app
- Slow boot

### Container

- Lightweight (MBs)

# Scaling & Load Balancing

One server isn't enough.

**Horizontal Scaling:** Add more containers.

**Load Balancer (Nginx):** Distributes traffic.

```
graph LR User --> LB[Load Balancer]; LB --> S1[Model Replica 1]; LB --> S2[Model Replica 2]; LB --> S3[Model Replica 3];
```

**Kubernetes (K8s):**

- Orchestrates containers.
- Auto-scaling: "If CPU > 80%, add replica".
- Self-healing: "If replica crashes, restart it".

# Deployment Strategies

## 1. Blue/Green Deployment:

- Run two environments: Blue (Current), Green (New).
- Switch traffic 100% to Green when ready.
- **Pros:** Instant rollback. **Cons:** 2x cost.

## 2. Canary Deployment:

- Send 10% traffic to V2, 90% to V1.
- Monitor errors. Gradually increase to 100%.
- **Pros:** Safer. **Cons:** Complex routing.

## 3. A/B Testing:

- Split traffic to measure *business impact* (not just errors).

# Model Drift: The Silent Killer

Code doesn't change, but **data** does.

1. **Data Drift:** Input distribution changes ( $P(X)$ ).

- *Example:* Training images were sunny, now users upload night photos.

2. **Concept Drift:** Relationship changes ( $P(Y|X)$ ).

- *Example:* "Spam" definition changes over time.

```
graph TD A[Training Data Dist] -- Time Passes --> B[Production Data Dist]; B --> C{Diff > Threshold?}; C -- Yes --> D[Alert & Retrain]; C -- No --> E[Keep Serving];
```

**Detection:**

- Statistical tests: Kolmogorov-Smirnov (KS) test.
- Tools: Evidently AI, Alibi Detect.

# API Design with FastAPI

**Pydantic** ensures inputs match your schema.

```
from fastapi import FastAPI
from pydantic import BaseModel, conlist

app = FastAPI()

# Define constraints
class InputData(BaseModel):
    # List of exactly 10 floats
    features: conlist(float, min_items=10, max_items=10)

@app.post("/predict")
def predict(data: InputData):
    # No need to check len(data.features), FastAPI did it!
    prediction = model.predict([data.features])
    return {"class": int(prediction[0])}
```

# Lab Overview: From Local to Cloud

Today's Lab:

1. **Serialize**: Save a scikit-learn model.
2. **API**: Wrap it in FastAPI.
3. **Containerize**: Write a Dockerfile.
4. **Optimize**: Convert to ONNX and benchmark speedup.
5. **Deploy**: (Optional) Push to Render/Heroku.

# Questions?

Let's ship some code!