RAG & Vector Databases

CS 203: Software Tools and Techniques for AI

Prof. Nipun Batra, IIT Gandhinagar

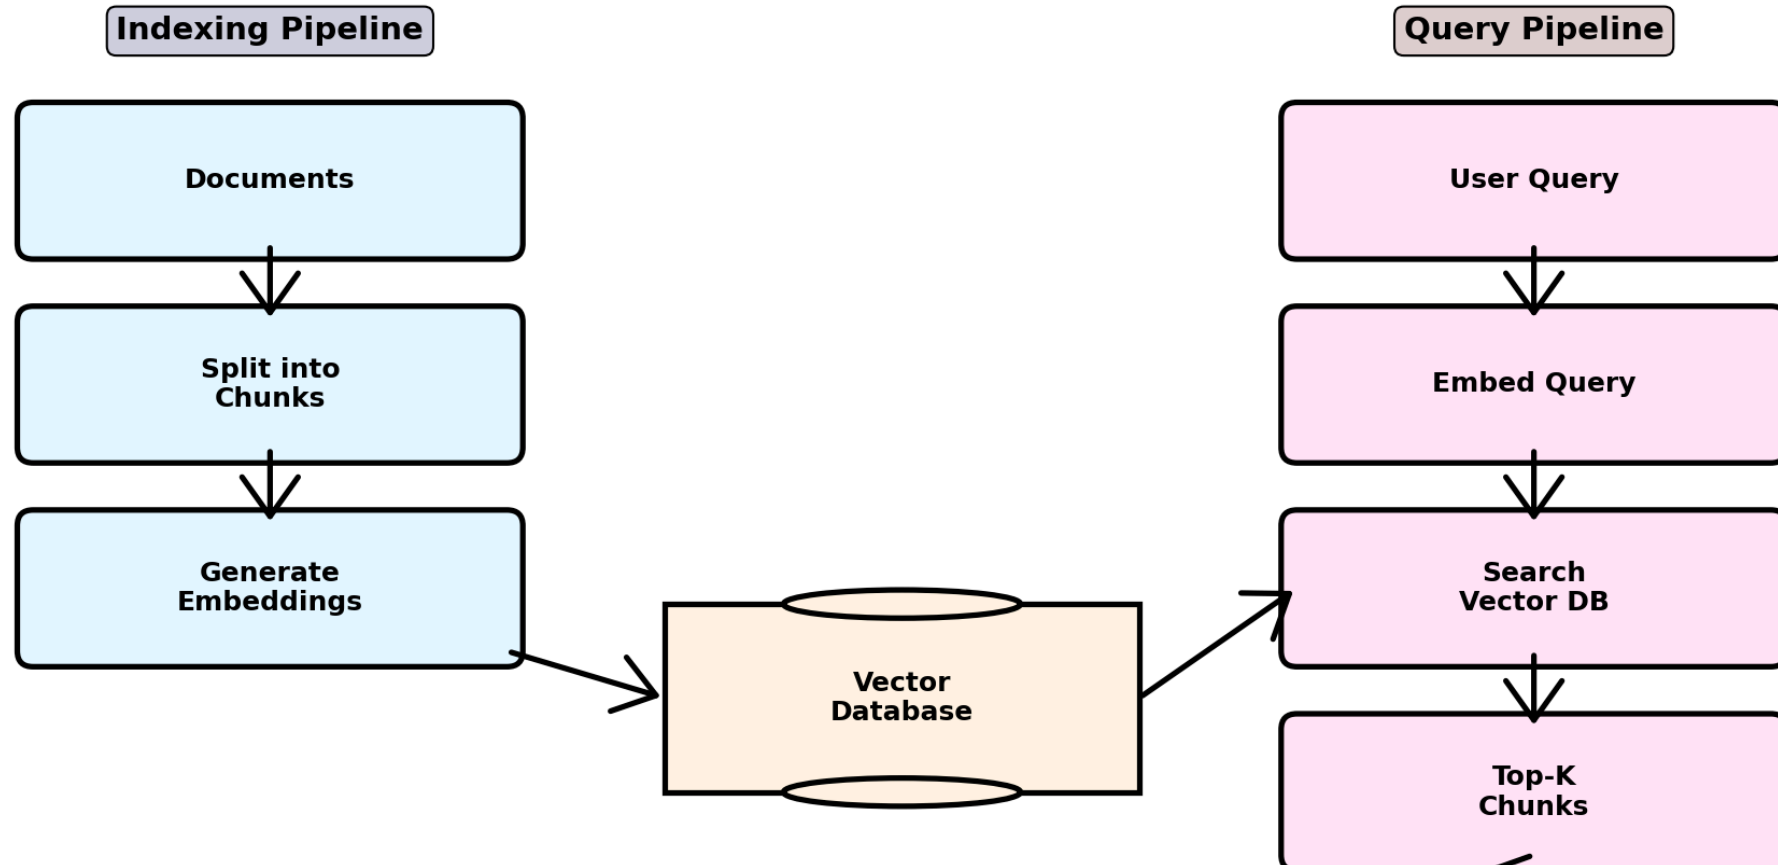# The LLM Knowledge Gap

LLMs are frozen in time.

- Trained on data up to a cutoff date (e.g., 2023).

- Don't know your private data (company emails, course syllabus).

- Can hallucinate facts.

Solution: Retrieval Augmented Generation (RAG)

- **Retrieve** relevant context from external source.

- **Augment** the prompt with this context.

- **Generate** answer using the augmented prompt.

RAG Architecture

RAG Architecture: Retrieval-Augmented Generation

Indexing Pipeline

Query Pipeline

Documents

Split into
Chunks

Generate
Embeddings

Vector
Database

User Query

Embed Query

Search
Vector DB

Top-K
Chunks

Embeddings: The Core Engine

What is an embedding?

- A vector (list of numbers) representing the *semantic meaning* of text.

- Similar texts have vectors close together in vector space.

Models:

- OpenAI `text-embedding-3-small` (1536 dim)

- Google `embedding-001`

- Open Source: `all-MiniLM-L6-v2` (Hugging Face)

```python
from sentence_transformers import SentenceTransformer
model = SentenceTransformer('all-MiniLM-L6-v2')

emb1 = model.encode("The cat sits outside")
emb2 = model.encode("A man is playing guitar")
```

Vector Mathematics: Embeddings as Points

Embeddings are vectors in high-dimensional space:

For a sentence, the embedding model produces:

$$\mathbf{v} = [v_1, v_2, \ldots, v_d] \in \mathbb{R}^d$$

where d is the dimensionality (e.g., 384, 768, 1536).

Example (384-dimensional embedding):

```
emb = model.encode("Hello world")
print(emb.shape)   # (384,)
print(emb[:5])     # [ 0.023, -0.145,  0.891, -0.234,  0.567]
```

Intuition: Each dimension captures a semantic feature.

- Similar words → similar coordinates

Similarity Metric 1: Cosine Similarity

Most common metric for embeddings.

Definition

$$\text{cosine\_sim}(\mathbf{a}, \mathbf{b}) = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\|\|\mathbf{b}\|} = \frac{\sum_{i=1}^{d} a_i b_i}{\sqrt{\sum_{i=1}^{d} a_i^2} \sqrt{\sum_{i=1}^{d} b_i^2}}$$

Range: $[-1, 1]$

- 1 = identical direction
- 0 = orthogonal (unrelated)
- -1 = opposite direction

Why cosine? Embeddings are normalized, so we care about direction, not magnitude.

Cosine Similarity: Worked Example

Given two embeddings (simplified to 3D):

- $\mathbf{a} = [1, 2, 3]$ (embedding for "cat")
- $\mathbf{b} = [2, 4, 6]$ (embedding for "feline")

Step 1: Dot Product

$$\mathbf{a} \cdot \mathbf{b} = (1)(2) + (2)(4) + (3)(6) = 2 + 8 + 18 = 28$$

Step 2: Magnitudes

$$\|\mathbf{a}\| = \sqrt{1^2 + 2^2 + 3^2} = \sqrt{14} \approx 3.742$$

$$\|\mathbf{b}\| = \sqrt{2^2 + 4^2 + 6^2} = \sqrt{56} \approx 7.483$$

Step 3: Cosine Similarity

Cosine Similarity in Python

```python
import numpy as np
from numpy.linalg import norm

def cosine_similarity(a, b):
    return np.dot(a, b) / (norm(a) * norm(b))

# Example
emb1 = model.encode("The cat sits outside")
emb2 = model.encode("A feline rests outdoors")
emb3 = model.encode("A man plays guitar")

print(f"cat vs feline: {cosine_similarity(emb1, emb2):.3f}")  # ~0.85
print(f"cat vs guitar: {cosine_similarity(emb1, emb3):.3f}")  # ~0.12
```

Interpretation:

- High similarity (> 0.7) → semantically similar
- Low similarity (< 0.3) → semantically different

Similarity Metric 2: Euclidean Distance

Measures straight-line distance in vector space.

Definition

$$d(\mathbf{a}, \mathbf{b}) = \|\mathbf{a} - \mathbf{b}\| = \sqrt{\sum_{i=1}^{d}(a_i - b_i)^2}$$

Range: $[0, \infty)$

- 0 = identical vectors
- Large value = far apart

Example (3D vectors):

- $\mathbf{a} = [1, 2, 3]$

Similarity Metric 3: Dot Product

Simplest metric (used when vectors are normalized).

Definition

$$\mathbf{a} \cdot \mathbf{b} = \sum_{i=1}^{d} a_i b_i$$

Range: $[-\infty, \infty]$ (or $[-1, 1]$ if normalized)

When to use:

- If embeddings are **already L2-normalized** (length = 1)
- Cosine similarity = dot product for normalized vectors
- Faster to compute (no division needed)

Comparison: Cosine vs Euclidean vs Dot Product

Metric	Formula	Range	Best For
Cosine	$\frac{\mathbf{a}\cdot\mathbf{b}}{|\mathbf{a}||\mathbf{b}|}$	$[-1, 1]$	Semantic similarity (direction)
Euclidean	$|\mathbf{a} - \mathbf{b}|$	$[0, \infty)$	Absolute distance
Dot Product	$\mathbf{a} \cdot \mathbf{b}$	$\mathbb{R}$	Normalized embeddings

Which to use for RAG?

- **Cosine** (default for most embedding models)

- Dot product if embeddings are pre-normalized (faster)

In practice: Most vector DBs (ChromaDB, Pinecone) default to cosine or dot.

Approximate Nearest Neighbors (ANN)

Problem: Finding exact nearest neighbors is $O(N \cdot d)$ (slow for millions of vectors).

Solution: Use approximate algorithms that trade accuracy for speed.

ANN Algorithms

1. **HNSW** (Hierarchical Navigable Small Worlds)

   - Graph-based search

   - Fast queries, high recall

   - Used by Pinecone, Qdrant

2. **IVF** (Inverted File Index)

   - Cluster vectors, search only relevant clusters

Chunking Strategies

Why chunk? LLMs have finite context windows, and retrieval is more precise with smaller chunks.

Common Strategies

1. **Fixed-size chunks**:

   - Split every 500 characters

   - Simple but can break mid-sentence

2. **Recursive Character Splitter** (LangChain default):

   - Try splitting by paragraph, then sentence, then character

   - Keeps semantic units together

3. **Semantic chunking**:

Chunking: Mathematical Perspective

Trade-off: Chunk size vs retrieval granularity.

Let:

- L = document length (tokens)
- C = chunk size
- O = overlap size

Number of chunks:

$$N_{\text{chunks}} = \left\lceil \frac{L - O}{C - O} \right\rceil$$

Example:

- Document: 10,000 tokens

Retrieval Evaluation Metrics

How do we know if our RAG system retrieves the right documents?

Recall@K

Definition: Of all relevant documents, what fraction appears in top-K results?

$$\text{Recall@K} = \frac{|\{\text{relevant docs}\} \cap \{\text{top-K results}\}|}{|\{\text{relevant docs}\}|}$$

Example:

- 5 relevant documents total
- Top-3 results contain 2 of them

$$\text{Recall@3} = \frac{2}{5} = 0.4$$

Mean Reciprocal Rank (MRR)

Measures how high the first relevant result ranks.

$$\text{MRR} = \frac{1}{|Q|} \sum_{i=1}^{|Q|} \frac{1}{\text{rank}_i}$$

where rank_i is the position of the first relevant result for query i.

Example:

- Query 1: First relevant doc at position 1 $\rightarrow 1/1 = 1.0$

- Query 2: First relevant doc at position 3 $\rightarrow 1/3 = 0.333$

- Query 3: First relevant doc at position 2 $\rightarrow 1/2 = 0.5$

$$\text{MRR} = \frac{1}{3}(1.0 + 0.333 + 0.5) = 0.611$$

Normalized Discounted Cumulative Gain (NDCG)

Accounts for both relevance and ranking position.

DCG@K (Discounted Cumulative Gain)

$$\text{DCG@K} = \sum_{i=1}^{K} \frac{\text{rel}_i}{\log_2(i+1)}$$

where rel_i is the relevance score of result at position i (e.g., 0 or 1, or graded).

NDCG@K (Normalized DCG)

$$\text{NDCG@K} = \frac{\text{DCG@K}}{\text{IDCG@K}}$$

where IDCG = DCG of the ideal ranking (all relevant docs first).

Vector Databases

Specialized databases for storing and searching high-dimensional vectors.

Why not standard SQL?

- SQL is good for exact match (`WHERE id = 5`).
- Vector DB is good for approximate nearest neighbor (ANN) search.

Popular Tools:

- **ChromaDB**: Open-source, local/in-memory (Great for dev).
- **Pinecone**: Managed service (Scalable).
- **FAISS**: Facebook's library for dense retrieval (The engine behind many DBs).
- **Qdrant**: Rust-based, fast.
- **pgvector**: Postgres extension.

ChromaDB Example

```python
import chromadb

# 1. Initialize Client
client = chromadb.Client()
collection = client.create_collection("course_docs")

# 2. Add Documents (Chroma handles embedding by default if not provided)
collection.add(
    documents=["CS203 covers AI tools.", "The exam is on Monday.", "Python is used."],
    metadatas=[{"source": "syllabus"}, {"source": "schedule"}, {"source": "intro"}],
    ids=["id1", "id2", "id3"]
)

# 3. Query
results = collection.query(
    query_texts=["When is the test?"],
    n_results=1
)

print(results['documents'])
# Output: [['The exam is on Monday.']]
```

Building a RAG Pipeline: Step 1 (Ingestion)

Chunking Matters: LLMs have context limits, and we want precise retrieval.

- Split by character count?

- Split by paragraph?

- Recursive character text splitter (LangChain).

```python
from langchain.text_splitter import RecursiveCharacterTextSplitter

text = "Long document..."
splitter = RecursiveCharacterTextSplitter(
    chunk_size=500,
    chunk_overlap=50
)
chunks = splitter.split_text(text)
# Now embed and store 'chunks'
```

Building a RAG Pipeline: Step 2 (Retrieval)

```python
# User asks: "How do I install the tools?"
query_vector = embedding_model.encode("How do I install the tools?")

# Search Vector DB
results = collection.query(query_embeddings=[query_vector], n_results=3)
context_text = "\n".join(results['documents'][0])
```

Building a RAG Pipeline: Step 3 (Generation)

```python
import google.generativeai as genai

prompt = f"""
You are a helpful teaching assistant. Answer the question based ONLY on the context below.

Context:
{context_text}

Question:
How do I install the tools?
"""

model = genai.GenerativeModel('gemini-pro')
response = model.generate_content(prompt)
print(response.text)
```

Orchestration Frameworks

Writing all this glue code is tedious. Frameworks help:

LangChain:

- Massive ecosystem.

- Chains, Agents, Integrations.

- Can be complex/verbose.

LlamaIndex:

- Specialized for data ingestion/retrieval.

- better for complex data structures (hierarchical indices).

Haystack:

- Pipeline-centric, robust.

LangChain Example

```python
from langchain.vectorstores import Chroma
from langchain.embeddings import OpenAIEmbeddings
from langchain.chains import RetrievalQA
from langchain.llms import OpenAI

# Setup
db = Chroma(persist_directory="./db", embedding_function=OpenAIEmbeddings())
retriever = db.as_retriever()
llm = OpenAI()

# Chain
qa = RetrievalQA.from_chain_type(
    llm=llm,
    chain_type="stuff",
    retriever=retriever
)

# Run
print(qa.run("What is the grading policy?"))
```

Advanced RAG Techniques

1. **Hybrid Search**:

   - Combine Vector Search (semantic) + Keyword Search (BM25).

   - Good for exact terms like product IDs or names.

2. **Re-ranking**:

   - Retrieve 50 docs quickly (Vector DB).

   - Re-rank top 50 using a slower, more accurate model (Cross-Encoder).

   - Pass top 5 to LLM.

3. **Query Expansion**:

   - LLM rewrites user query into multiple versions.

   - Search all versions, deduplicate results.

Lab: Chat with Your PDF

Objective: Build a tool to upload a PDF and ask questions about it.

Tools:

- **pypdf**: Extract text.
- **RecursiveCharacterTextSplitter**: Chunking.
- **ChromaDB**: Vector Store.
- **Gemini/OpenAI API**: Embeddings & Generation.
- **Streamlit**: UI.

Workflow:

1. User uploads `paper.pdf`.
2. App extracts text -> chunks -> embeds -> stores in session ChromaDB.

Resources

- **Pinecone Learning Center**: pinecone.io/learn

- **LangChain Docs**: python.langchain.com

- **ChromaDB**: trychroma.com

- **DeepLearning.AI Short Courses**: "Building Systems with LLM API"

Questions?